

loan-approval-prediction-3

October 31, 2024

```
[3]: # Importing necessary libraries
import pandas as pd
import numpy as np
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import classification_report, roc_auc_score, accuracy_score
from imblearn.over_sampling import SMOTE
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
import xgboost as xgb
from imblearn.under_sampling import RandomUnderSampler
from sklearn.model_selection import GridSearchCV
import matplotlib.pyplot as plt
import plotly.express as px
import plotly.subplots as sp
import plotly.graph_objects as go
from plotly.subplots import make_subplots
```

```
# Suppress warnings for clarity
import warnings
warnings.filterwarnings("ignore")
# Set Seaborn color palette for the notebook
sns.set_palette("deep")
sns.set_theme(style="whitegrid", palette="deep")
```

```
[4]: # Load Train Data
df_train = pd.read_csv('loan_train.csv')
df_train
```

```
[4]:
```

	id	person_age	person_income	person_home_ownership	\
0	0	37	35000	RENT	
1	1	22	56000	OWN	
2	2	29	28800	OWN	

3	3	30	70000	RENT
4	4	22	60000	RENT
...
58640	58640	34	120000	MORTGAGE
58641	58641	28	28800	RENT
58642	58642	23	44000	RENT
58643	58643	22	30000	RENT
58644	58644	31	75000	MORTGAGE

	person_emp_length	loan_intent	loan_grade	loan_amnt	loan_int_rate	\
0	0.0	EDUCATION	B	6000	11.49	
1	6.0	MEDICAL	C	4000	13.35	
2	8.0	PERSONAL	A	6000	8.90	
3	14.0	VENTURE	B	12000	11.11	
4	2.0	MEDICAL	A	6000	6.92	
...	
58640	5.0	EDUCATION	D	25000	15.95	
58641	0.0	MEDICAL	C	10000	12.73	
58642	7.0	EDUCATION	D	6800	16.00	
58643	2.0	EDUCATION	A	5000	8.90	
58644	2.0	VENTURE	B	15000	11.11	

	loan_percent_income	cb_person_default_on_file	\
0	0.17	N	
1	0.07	N	
2	0.21	N	
3	0.17	N	
4	0.10	N	
...	
58640	0.21	Y	
58641	0.35	N	
58642	0.15	N	
58643	0.17	N	
58644	0.20	N	

	cb_person_cred_hist_length	loan_status
0	14	0
1	2	0
2	10	0
3	5	0
4	3	0
...
58640	10	0
58641	8	1
58642	2	1
58643	3	0
58644	5	0

[58645 rows x 13 columns]

```
[5]: # Load Test Data
df_test = pd.read_csv('loan_test.csv')
df_test
```

```
[5]:      id  person_age  person_income  person_home_ownership  \
0    58645         23         69000             RENT
1    58646         26         96000             MORTGAGE
2    58647         26         30000             RENT
3    58648         33         50000             RENT
4    58649         26        102000             MORTGAGE
...    ...         ...         ...             ...
39093  97738         22         31200             MORTGAGE
39094  97739         22         48000             MORTGAGE
39095  97740         51         60000             MORTGAGE
39096  97741         22         36000             MORTGAGE
39097  97742         31         45000             RENT

      person_emp_length      loan_intent  loan_grade  loan_amnt  \
0                3.0    HOMEIMPROVEMENT            F    25000
1                6.0            PERSONAL            C    10000
2                5.0            VENTURE            E     4000
3                4.0    DEBTCONSOLIDATION            A     7000
4                8.0    HOMEIMPROVEMENT            D    15000
...                ...                ...            ...
39093            2.0    DEBTCONSOLIDATION            B     3000
39094            6.0            EDUCATION            A     7000
39095            0.0            PERSONAL            A    15000
39096            4.0            PERSONAL            D    14000
39097            6.0    DEBTCONSOLIDATION            B    19450

      loan_int_rate  loan_percent_income  cb_person_default_on_file  \
0             15.76                0.36                        N
1             12.68                0.10                        Y
2             17.19                0.13                        Y
3              8.90                0.14                        N
4             16.32                0.15                        Y
...                ...                ...
39093            10.37                0.10                        N
39094             6.03                0.15                        N
39095             7.51                0.25                        N
39096            15.62                0.39                        Y
39097             9.91                0.44                        N

      cb_person_cred_hist_length
```

0	2
1	4
2	2
3	7
4	4
...	...
39093	4
39094	3
39095	25
39096	4
39097	9

[39098 rows x 12 columns]

```
[6]: # Data Overview
df_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 58645 entries, 0 to 58644
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                     58645 non-null  int64
1   person_age                           58645 non-null  int64
2   person_income                        58645 non-null  int64
3   person_home_ownership                58645 non-null  object
4   person_emp_length                    58645 non-null  float64
5   loan_intent                           58645 non-null  object
6   loan_grade                           58645 non-null  object
7   loan_amnt                            58645 non-null  int64
8   loan_int_rate                        58645 non-null  float64
9   loan_percent_income                  58645 non-null  float64
10  cb_person_default_on_file             58645 non-null  object
11  cb_person_cred_hist_length            58645 non-null  int64
12  loan_status                           58645 non-null  int64
dtypes: float64(3), int64(6), object(4)
memory usage: 5.8+ MB
```

```
[7]: df_test.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 39098 entries, 0 to 39097
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                     39098 non-null  int64
1   person_age                           39098 non-null  int64
2   person_income                        39098 non-null  int64
```

```

3  person_home_ownership      39098 non-null object
4  person_emp_length          39098 non-null float64
5  loan_intent                 39098 non-null object
6  loan_grade                 39098 non-null object
7  loan_amnt                  39098 non-null int64
8  loan_int_rate              39098 non-null float64
9  loan_percent_income        39098 non-null float64
10 cb_person_default_on_file  39098 non-null object
11 cb_person_cred_hist_length 39098 non-null int64
dtypes: float64(3), int64(5), object(4)
memory usage: 3.6+ MB

```

```
[8]: df_train.describe()
```

```

[8]:
      count      id      person_age  person_income  person_emp_length  \
count  58645.000000  58645.000000  5.864500e+04  58645.000000
mean    29322.000000    27.550857  6.404617e+04    4.701015
std     16929.497605     6.033216  3.793111e+04    3.959784
min         0.000000    20.000000  4.200000e+03    0.000000
25%    14661.000000    23.000000  4.200000e+04    2.000000
50%    29322.000000    26.000000  5.800000e+04    4.000000
75%    43983.000000    30.000000  7.560000e+04    7.000000
max    58644.000000   123.000000  1.900000e+06   123.000000

      count      loan_amnt  loan_int_rate  loan_percent_income  \
count  58645.000000  58645.000000  58645.000000
mean     9217.556518    10.677874    0.159238
std     5563.807384     3.034697    0.091692
min       500.000000     5.420000    0.000000
25%      5000.000000     7.880000    0.090000
50%      8000.000000    10.750000    0.140000
75%     12000.000000    12.990000    0.210000
max     35000.000000    23.220000    0.830000

      count      cb_person_cred_hist_length  loan_status
count    58645.000000  58645.000000  58645.000000
mean         5.813556    0.142382
std         4.029196    0.349445
min          2.000000    0.000000
25%          3.000000    0.000000
50%          4.000000    0.000000
75%          8.000000    0.000000
max         30.000000    1.000000

```

```
[9]: df_test.describe()
```

```
[9]:
```

	id	person_age	person_income	person_emp_length	\
count	39098.000000	39098.000000	3.909800e+04	39098.000000	
mean	78193.500000	27.566781	6.406046e+04	4.687068	
std	11286.764749	6.032761	3.795583e+04	3.868395	
min	58645.000000	20.000000	4.000000e+03	0.000000	
25%	68419.250000	23.000000	4.200000e+04	2.000000	
50%	78193.500000	26.000000	5.800000e+04	4.000000	
75%	87967.750000	30.000000	7.588500e+04	7.000000	
max	97742.000000	94.000000	1.900000e+06	42.000000	

	loan_amnt	loan_int_rate	loan_percent_income	\
count	39098.000000	39098.000000	39098.000000	
mean	9251.466188	10.661216	0.159573	
std	5576.254680	3.020220	0.091633	
min	700.000000	5.420000	0.000000	
25%	5000.000000	7.880000	0.090000	
50%	8000.000000	10.750000	0.140000	
75%	12000.000000	12.990000	0.210000	
max	35000.000000	22.110000	0.730000	

	cb_person_cred_hist_length
count	39098.000000
mean	5.830707
std	4.072157
min	2.000000
25%	3.000000
50%	4.000000
75%	8.000000
max	30.000000

```
[10]: # Train and Test Data Missing Values
print(df_train.isnull().sum())
print(df_test.isnull().sum())
```

id	0
person_age	0
person_income	0
person_home_ownership	0
person_emp_length	0
loan_intent	0
loan_grade	0
loan_amnt	0
loan_int_rate	0
loan_percent_income	0
cb_person_default_on_file	0
cb_person_cred_hist_length	0
loan_status	0

```

dtype: int64
id                0
person_age        0
person_income     0
person_home_ownership  0
person_emp_length  0
loan_intent        0
loan_grade         0
loan_amnt          0
loan_int_rate      0
loan_percent_income  0
cb_person_default_on_file  0
cb_person_cred_hist_length  0
dtype: int64

```

```

[11]: # Duplicate Check
      duplicates = df_train.duplicated()
      print(df_train[duplicates])

```

```

Empty DataFrame
Columns: [id, person_age, person_income, person_home_ownership,
person_emp_length, loan_intent, loan_grade, loan_amnt, loan_int_rate,
loan_percent_income, cb_person_default_on_file, cb_person_cred_hist_length,
loan_status]
Index: []

```

```

[12]: duplicates_test = df_test.duplicated()
      print(df_test[duplicates_test])

```

```

Empty DataFrame
Columns: [id, person_age, person_income, person_home_ownership,
person_emp_length, loan_intent, loan_grade, loan_amnt, loan_int_rate,
loan_percent_income, cb_person_default_on_file, cb_person_cred_hist_length]
Index: []

```

```

[13]: import matplotlib.pyplot as plt
      import seaborn as sns

      # Count the values in loan_status column (1s and 0s)
      loan_status_counts = df_train['loan_status'].value_counts()

      # Set figure size
      plt.figure(figsize=(7, 5))

      # Create a barplot using seaborn
      sns.barplot(x=loan_status_counts.index,
                  y=loan_status_counts.values,
                  palette="viridis")

```

```

# Add labels for each bar
for index, value in enumerate(loan_status_counts.values):
    plt.text(index, value, f'{value}', ha='center', va='bottom', fontsize=10)

# Titles and labels
plt.title('Count of 1 and 0 in Loan Status', fontsize=16)
plt.xlabel('Loan Status (1 = Approved, 0 = Not Approved)', fontsize=12)
plt.ylabel('Count', fontsize=12)

# Show the plot
plt.show()


```

C:\Users\Dell\AppData\Local\Temp\ipykernel_10416\3937357170.py:11:

FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x=loan_status_counts.index,
```




```
[23]: # Import libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import plotly.express as px
import plotly.subplots as sp
import plotly.graph_objects as go
from plotly.subplots import make_subplots
# Importing essential libraries
from sklearn.model_selection import StratifiedKFold
from sklearn.metrics import roc_auc_score
from sklearn.preprocessing import LabelEncoder
import lightgbm as lgb
import xgboost as xgb
from catboost import CatBoostClassifier
from sklearn.ensemble import VotingClassifier
from IPython.display import display, HTML

from sklearn.metrics import roc_auc_score
from sklearn.model_selection import StratifiedKFold
import random

# Suppress warnings for clarity
import warnings
warnings.filterwarnings("ignore")
# Set Seaborn color palette for the notebook
sns.set_palette("deep")
sns.set_theme(style="whitegrid", palette="deep")

[25]: # Preparing data for the donut chart (loan intent by age)
loan_intent_age_distribution = df_train.groupby(['person_age', 'loan_intent']).
    ↪size().reset_index(name='count')

# Create a combined string for hover data
loan_intent_age_distribution['hover_data'] = (
    loan_intent_age_distribution['person_age'].astype(str) +
    " - " + loan_intent_age_distribution['loan_intent']
)

# Preparing data for the second donut chart (overall loan intent distribution)
overall_loan_intent_distribution = df_train.groupby('loan_intent').size().
    ↪reset_index(name='count')

# Define a deep color palette
deep_colors = ['#1f77b4', '#ff7f0e', '#2ca02c', '#d62728', '#9467bd',
    ↪'#8c564b', '#e377c2', '#7f7f7f', '#bcbd22', '#17becf']
```

```

# Create the donut chart for loan intent by age
donut_chart_age = go.Pie(
    labels=loan_intent_age_distribution['loan_intent'],
    values=loan_intent_age_distribution['count'],
    name="Loan Intent by Age",
    hole=0.4, # Set hole size for donut chart
    marker=dict(
        line=dict(
            color='rgba(0, 0, 0, 0.2)', # Shadow color
            width=4 # Width of the shadow effect
        ),
        colors=deep_colors # Use the deep color palette
    ),
    hovertemplate=(
        "<b>Details:</b> %{customdata}<br>" # Display combined age and loan
        ↪intent
        "<b>Count:</b> %{value}<extra></extra>" # Display count without extra
        ↪hover box
    ),
    customdata=loan_intent_age_distribution['hover_data'].values # Use
    ↪combined data for custom hover
)

# Create the donut chart for overall loan intent distribution
donut_chart_overall = go.Pie(
    labels=overall_loan_intent_distribution['loan_intent'],
    values=overall_loan_intent_distribution['count'],
    name="Overall Loan Intent Distribution",
    hole=0.4, # Set hole size for donut chart
    marker=dict(
        line=dict(
            color='rgba(0, 0, 0, 0.2)', # Shadow color
            width=4 # Width of the shadow effect
        ),
        colors=deep_colors # Use the same deep color palette
    ),
    hovertemplate=(
        "<b>Loan Intent:</b> %{label}<br>" # Display loan intent
        "<b>Count:</b> %{value}<extra></extra>" # Display count without extra
        ↪hover box
    )
)

# Create subplots
fig = make_subplots(rows=1, cols=2, specs=[[{'type': 'pie'}, {'type': 'pie'}]],

```

```

        subplot_titles=("Loan Intent Distribution by Age", "Overall Loan Intent Distribution"))


# Add donut charts to the figure
fig.add_trace(donut_chart_age, row=1, col=1)
fig.add_trace(donut_chart_overall, row=1, col=2)

# Step 6: Update layout to increase figure size and position the legend
fig.update_layout(
    height=500, width=900, # Set figure size
    showlegend=True,
    legend=dict(
        x=1, # Place the legend outside
        y=1,
        traceorder="normal",
        font=dict(size=12)
    ),
    title="Loan Intent Insights"
)

# Display the figure
fig.show()

```

Loan Intent Insights



```

[27]: # Get counts of each category
home_ownership_counts = df_train['person_home_ownership'].value_counts().
    ↪reset_index()
home_ownership_counts.columns = ['person_home_ownership', 'count']

# Define a deep color palette

```

```

deep_color_palette = ['#1f77b4', '#ff7f0e', '#2ca02c', '#d62728', '#9467bd',
↳ '#8c564b', '#e377c2'] # Add more colors if needed

# Create subplots with two columns, specifying the second column as a pie chart
fig = make_subplots(rows=1, cols=2,
                    specs=[[{"type": "bar"}, {"type": "pie"}]],
                    subplot_titles=('Bar Plot', 'Pie Chart'))


# Add the bar trace
fig.add_trace(go.Bar(
    x=home_ownership_counts['person_home_ownership'],
    y=home_ownership_counts['count'],
    text=home_ownership_counts['count'],
    textposition='auto',
    marker=dict(color=deep_color_palette[:len(home_ownership_counts)]), # Use
↳ the deep color palette
    hovertemplate='Home Ownership: %{x}<br>Count: %{y}<extra></extra>' #
↳ Display column names on hover
), row=1, col=1) # Bar plot in the first subplot

# Add the pie chart
fig.add_trace(go.Pie(
    labels=home_ownership_counts['person_home_ownership'],
    values=home_ownership_counts['count'],
    marker=dict(colors=deep_color_palette[:len(home_ownership_counts)]), # Use
↳ the deep color palette
    hoverinfo='label+percent+value',
    textinfo='label+value',
    hovertemplate='Home Ownership: %{label}<br>Count: %{value}<br>Percent:
↳ %{percent:.2%}<extra></extra>' # Display column names and percent on hover
), row=1, col=2) # Pie chart in the second subplot

# Update layout for better visualization
fig.update_layout(
    title='Home Ownership Distribution',
    height=400,
    width=800,
    yaxis=dict(
        tickformat=',g', # Use thousands separator without abbreviating
        showspikes=False # Disable spikes
    )
)

# Show the figure
fig.show()

```



```
[29]: fig = make_subplots(
    rows=1, cols=2,
    specs=[[{'type': 'pie'}, {'type': 'bar'}]],
    subplot_titles=("Loan Status Donut Chart", "Home Ownership Count Bar Plot")
)

# Prepare data for the donut chart (loan_status)
loan_status_counts = df_train['loan_status'].value_counts().reset_index()
loan_status_counts.columns = ['loan_status', 'count']

# Prepare the donut chart
donut_chart = go.Pie(
    labels=loan_status_counts['loan_status'],
    values=loan_status_counts['count'],
    hole=0.4, # Set hole size for donut chart
    marker=dict(line=dict(color='rgba(0, 0, 0, 0.2)', width=2)) # Add shadow
    ↪effect
)

# Add the donut chart to the figure
fig.add_trace(donut_chart, row=1, col=1)

# Prepare data for the bar plot (person_home_ownership)
home_ownership_counts = df_train['person_home_ownership'].value_counts().
    ↪reset_index()
home_ownership_counts.columns = ['person_home_ownership', 'count']

# Define a deep color palette
deep_color_palette = ['#1f77b4', '#ff7f0e', '#2ca02c', '#d62728', '#9467bd',
    ↪'#8c564b'] # More colors can be added if needed

# Add the bar trace
```

```


fig.add_trace(go.Bar(
    x=home_ownership_counts['person_home_ownership'],
    y=home_ownership_counts['count'],
    text=home_ownership_counts['count'], # Display counts above bars
    textposition='auto', # Automatically position the text
    marker=dict(color=deep_color_palette[:len(home_ownership_counts)]) # Use
    ↪ colors from the deep palette
), row=1, col=2)

# Update layout for better visualization
fig.update_layout(
    title='Loan Status and Home Ownership Analysis',
    height=400,
    width=900,
    showlegend=True,
    yaxis_tickformat=',', # Use a comma as a thousands separator
)

# Show the figure
fig.show()

```

Loan Status and Home Ownership Analysis



```

[31]: # Create subplots
fig = make_subplots(
    rows=1, cols=3,
    specs=[[{'type': 'bar'}, {'type': 'bar'}, {'type': 'bar'}]],
    subplot_titles=("Home Ownership Count", "Loan Intent Count", "Loan Grade
    ↪ Count")
)

# Prepare data for each column

# Count for person_home_ownership

```

```

home_ownership_counts = df_train['person_home_ownership'].value_counts().
    ↪reset_index()
home_ownership_counts.columns = ['person_home_ownership', 'count']

# Define a deep color for home ownership
home_ownership_color = '#FF4500' # OrangeRed

# Add the bar trace for home ownership
fig.add_trace(go.Bar(
    x=home_ownership_counts['person_home_ownership'],
    y=home_ownership_counts['count'],
    name='Home Ownership', # Legend name
    marker=dict(color=home_ownership_color),
    text=home_ownership_counts['count'], # Display counts above bars
    textposition='auto' # Automatically position the text
), row=1, col=1)

# Count for loan_intent
loan_intent_counts = df_train['loan_intent'].value_counts().reset_index()
loan_intent_counts.columns = ['loan_intent', 'count']

# Define a deep color for loan intent
loan_intent_color = '#1E90FF' # Dodger Blue

# Add the bar trace for loan intent
fig.add_trace(go.Bar(
    x=loan_intent_counts['loan_intent'],
    y=loan_intent_counts['count'],
    name='Loan Intent', # Legend name
    marker=dict(color=loan_intent_color),
    text=loan_intent_counts['count'], # Display counts above bars
    textposition='auto' # Automatically position the text
), row=1, col=2)

# Count for loan_grade
loan_grade_counts = df_train['loan_grade'].value_counts().reset_index()
loan_grade_counts.columns = ['loan_grade', 'count']

# Define a deep color for loan grade
loan_grade_color = '#8A2BE2' # BlueViolet

# Add the bar trace for loan grade
fig.add_trace(go.Bar(
    x=loan_grade_counts['loan_grade'],
    y=loan_grade_counts['count'],
    name='Loan Grade', # Legend name
    marker=dict(color=loan_grade_color),

```

```


    text=loan_grade_counts['count'], # Display counts above bars
    textposition='auto' # Automatically position the text
), row=1, col=3)

# Update layout for better visualization
fig.update_layout(
    title='Count Plots for Home Ownership, Loan Intent, and Loan Grade',
    height=400,
    width=1200,
    showlegend=True,
)

# Update y-axis properties to display full numbers
for i in range(1, 4):
    fig.update_yaxes(
        tickformat='g', # Use thousands separator without abbreviating
        showspikes=False, # Disable spikes on y-axis
        row=1, col=i # Specify which subplot to update
    )

# Show the figure
fig.show()

```



```

[33]: # Set the style for the plots
sns.set(style="whitegrid")

# Define a deeper color palette
colors = ['#FF4500', '#1E90FF', '#8A2BE2'] # Deep colors: OrangeRed, Dodger_
        ↪ Blue, BlueViolet

# Create a figure with subplots
fig, axs = plt.subplots(2, 2, figsize=(14, 10))

```



```

# Violin plot for person_age
sns.violinplot(x='loan_status', y='person_age', data=df_train, ax=axes[0, 0],
    palette=colors[:2], inner="quartile")
axes[0, 0].set_title('Violin Plot of Age by Loan Status')
axes[0, 0].set_ylabel('Person Age')

# Pie chart for loan status
loan_counts = df_train['loan_status'].value_counts()
axes[0, 1].pie(loan_counts, labels=loan_counts.index, autopct='%1.1f%%',
    startangle=90, colors=colors[:2])
axes[0, 1].set_title('Loan Status Distribution')

# Count plot for person_home_ownership by loan status
sns.countplot(x='person_home_ownership', hue='loan_status', data=df_train,
    ax=axes[1, 0], palette=colors)
axes[1, 0].set_title('Home Ownership by Loan Status')
axes[1, 0].set_ylabel('Counts')
axes[1, 0].set_xlabel('Home Ownership')

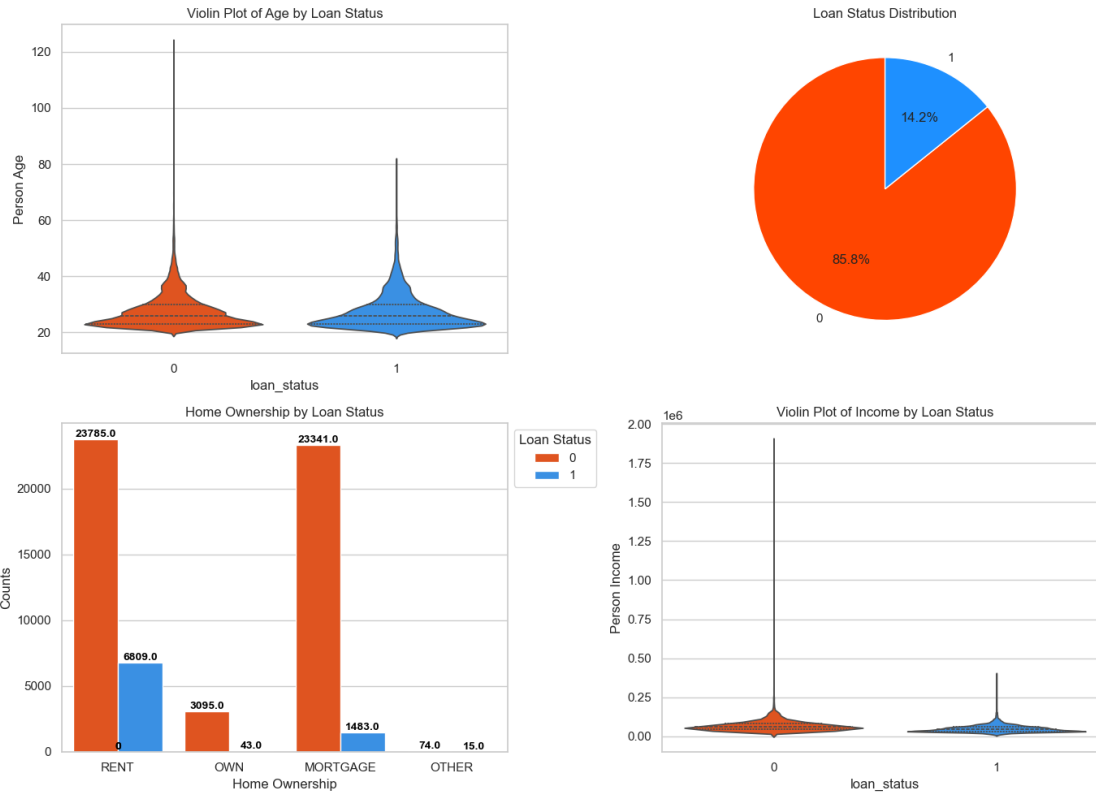
# Add value annotations for count plot
for p in axes[1, 0].patches:
    axes[1, 0].annotate(f'{p.get_height()}',
        (p.get_x() + p.get_width() / 2., p.get_height()),
        ha='center', va='bottom',
        color='black', size='small', weight='semibold')

# Violin plot for person_income
sns.violinplot(x='loan_status', y='person_income', data=df_train, ax=axes[1, 1],
    palette=colors[:2], inner="quartile")
axes[1, 1].set_title('Violin Plot of Income by Loan Status')
axes[1, 1].set_ylabel('Person Income')

# Adjust legend positions
axes[1, 0].legend(title='Loan Status', labels=loan_counts.index.tolist(),
    loc='upper left', bbox_to_anchor=(1, 1))

# Adjust layout
plt.tight_layout()
plt.show()

```



```
[35]: # Set the style for the plots
sns.set(style="whitegrid")

# Define a deeper color palette with rich tones
colors = ['#1F77B4', '#FF7F0E'] # Deep blue and deep orange

# Create a figure with subplots, increasing the figsize
fig, axs = plt.subplots(2, 2, figsize=(30, 20)) # Adjusted size for prominence

# Count plot for person_home_ownership by loan_status
sns.countplot(x='person_home_ownership', hue='loan_status', data=df_train,
              ↪ax=axs[0, 0], palette=colors)
axs[0, 0].set_title('Home Ownership by Loan Status', fontsize=20)
axs[0, 0].set_ylabel('Counts', fontsize=16)
axs[0, 0].set_xlabel('Home Ownership', fontsize=16)
axs[0, 0].tick_params(labelsize=14) # Adjust tick label size

# Add value annotations for person_home_ownership count plot
for p in axs[0, 0].patches:
    axs[0, 0].annotate(f'{p.get_height()}',
                      (p.get_x() + p.get_width() / 2., p.get_height()),
                      ha='center', va='bottom',
```

```

        color='black', size='medium')

# Count plot for loan_intent by loan_status
sns.countplot(x='loan_intent', hue='loan_status', data=df_train, ax=axes[0, 1],
    palette=colors)
axes[0, 1].set_title('Loan Intent by Loan Status', fontsize=20)
axes[0, 1].set_ylabel('Counts', fontsize=16)
axes[0, 1].set_xlabel('Loan Intent', fontsize=16)
axes[0, 1].tick_params(labelsize=14) # Adjust tick label size

# Add value annotations for loan_intent count plot
for p in axes[0, 1].patches:
    axes[0, 1].annotate(f'{p.get_height()}',
        (p.get_x() + p.get_width() / 2., p.get_height()),
        ha='center', va='bottom',
        color='black', size='medium')

# Count plot for loan_grade by loan_status
sns.countplot(x='loan_grade', hue='loan_status', data=df_train, ax=axes[1, 0],
    palette=colors)
axes[1, 0].set_title('Loan Grade by Loan Status', fontsize=20)
axes[1, 0].set_ylabel('Counts', fontsize=16)
axes[1, 0].set_xlabel('Loan Grade', fontsize=16)
axes[1, 0].tick_params(labelsize=14) # Adjust tick label size

# Add value annotations for loan_grade count plot
for p in axes[1, 0].patches:
    axes[1, 0].annotate(f'{p.get_height()}',
        (p.get_x() + p.get_width() / 2., p.get_height()),
        ha='center', va='bottom',
        color='black', size='medium')


# Count plot for cb_person_default_on_file by loan_status
sns.countplot(x='cb_person_default_on_file', hue='loan_status', data=df_train,
    ax=axes[1, 1], palette=colors)
axes[1, 1].set_title('Default on File by Loan Status', fontsize=20)
axes[1, 1].set_ylabel('Counts', fontsize=16)
axes[1, 1].set_xlabel('Default on File', fontsize=16)
axes[1, 1].tick_params(labelsize=14) # Adjust tick label size

# Add value annotations for cb_person_default_on_file count plot
for p in axes[1, 1].patches:
    axes[1, 1].annotate(f'{p.get_height()}',
        (p.get_x() + p.get_width() / 2., p.get_height()),
        ha='center', va='bottom',
        color='black', size='medium')

```

```
# Adjust legends
for ax in axs.flat:
    ax.legend(title='Loan Status', fontsize=18, title_fontsize=20)

# Adjust layout
plt.tight_layout()
plt.show()
```



```
[37]: # Create scatter plots with Plotly Express
scatter_age = px.scatter(df_train, x='person_age', y='loan_amnt',
    color='loan_grade',
    title='Loan Amount vs Age Colored by Loan Grade',
    labels={'person_age': 'Person Age', 'loan_amnt': 'Loan Amount'},
    color_discrete_sequence=px.colors.qualitative.Prism)

scatter_income = px.scatter(df_train, x='person_income', y='loan_amnt',
    color='loan_grade',
    title='Loan Amount vs Income Colored by Loan Grade',
    labels={'person_income': 'Person Income',
    'loan_amnt': 'Loan Amount'},
    color_discrete_sequence=px.colors.qualitative.Prism)
```

```

# Set up a subplot with 1 row and 2 columns
fig = sp.make_subplots(rows=1, cols=2, subplot_titles=('Loan Amount vs Age',
↳ 'Loan Amount vs Income'))

# Add the scatter plots to the subplot figure
# For the first subplot, we keep the legend
for trace in scatter_age['data']:
    fig.add_trace(trace, row=1, col=1)

# For the second subplot, we hide the legend by setting `showlegend=False`
for trace in scatter_income['data']:
    trace.showlegend = False # Disable legend for the second plot
    fig.add_trace(trace, row=1, col=2)

# Update layout for the overall figure
fig.update_layout(
    title_text="Scatter Plot Subplots of Loan Amount by Age and Income",
    showlegend=True, # Display legend only once
    legend_title="Loan Grade",
    height=600, width=1200
)


# Ensure y-axis values show full numbers (e.g., 10000 instead of 10k)
fig.update_yaxes(tickformat="," , title_text="Loan Amount", row=1, col=1) #
↳ First subplot
fig.update_yaxes(tickformat="," , title_text="Loan Amount", row=1, col=2) #
↳ Second subplot

# Ensure x-axis values show full numbers for the second subplot (Loan Amount vs
↳ Income)
fig.update_xaxes(tickformat="," , title_text="Person Income", row=1, col=2) #
↳ Second subplot

fig.show()

```

Scatter Plot Subplots of Loan Amount by Age and Income



```
[39]: # Define the columns for which you want to create KDE plots
columns = ['person_age', 'person_income', 'loan_amnt', '
↳ 'cb_person_cred_hist_length']

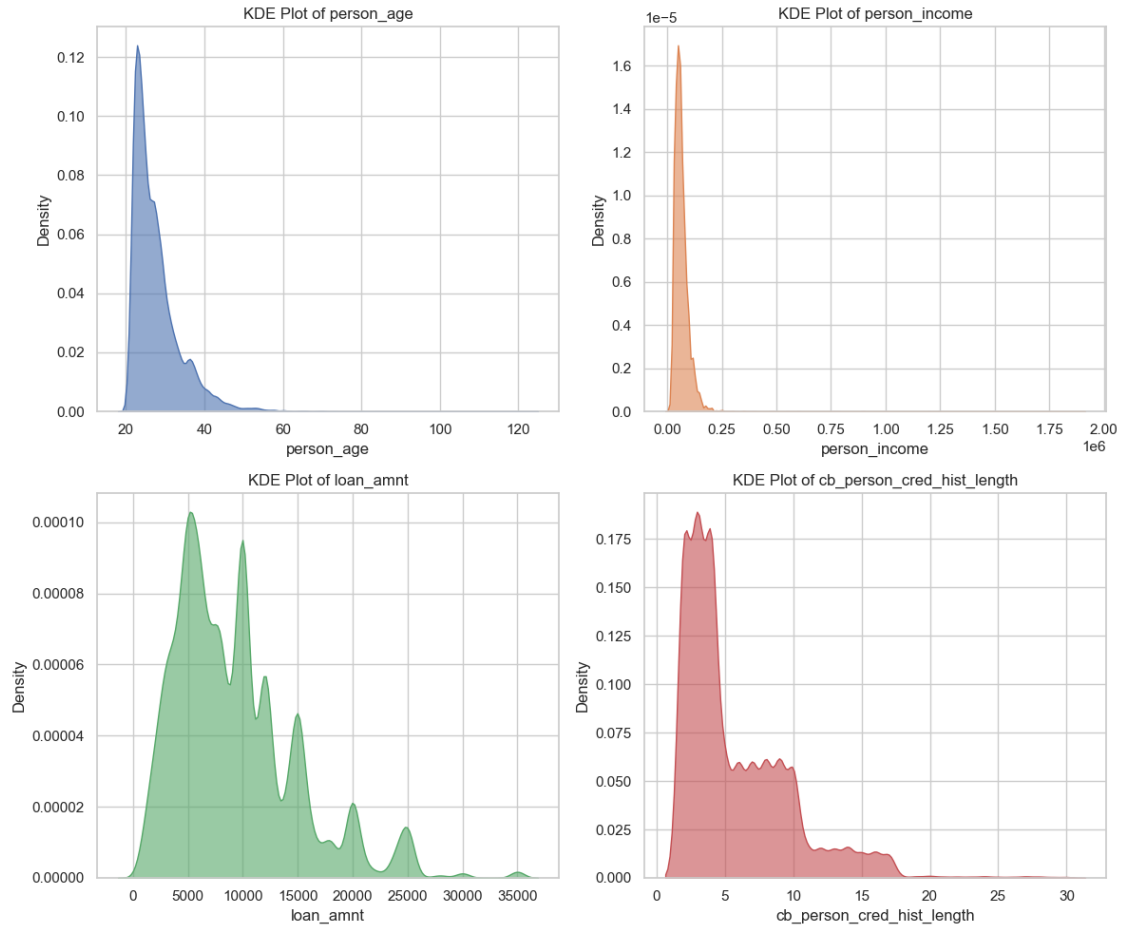
# Set a deep color palette
deep_palette = sns.color_palette("deep", len(columns))

# Create a figure with subplots
plt.figure(figsize=(12, 10))

# Loop through the columns and create a KDE plot for each
for i, column in enumerate(columns):
    plt.subplot(2, 2, i + 1) # Create a 2x2 grid of subplots
    sns.kdeplot(data=df_train[column], fill=True, color=deep_palette[i],
↳ alpha=0.6) # Create the KDE plot with a specific color
    plt.title(f'KDE Plot of {column}') # Set the title for each subplot
    plt.xlabel(column) # Label the x-axis
    plt.ylabel('Density') # Label the y-axis

# Adjust layout to prevent overlap
plt.tight_layout()

# Show the plots
plt.show()
```



```
[290]: # Create new features for df_train

# 1. Income to Loan Ratio
df_train['income_loan_ratio'] = df_train['person_income'] /_
    ↪df_train['loan_amnt']

# 2. Age-Employment Length Ratio
df_train['age_emp_length_ratio'] = df_train['person_emp_length'] /_
    ↪df_train['person_age']

# # 3. Loan to Employment Ratio
# df_train['loan_emp_ratio'] = df_train['loan_amnt'] /_
    ↪df_train['person_emp_length']

# 4. Loan Interest to Income Ratio
df_train['loan_int_income_ratio'] = df_train['loan_int_rate'] /_
    ↪df_train['person_income']
```

```

# 5. Credit History to Age Ratio
df_train['cred_hist_age_ratio'] = df_train['cb_person_cred_hist_length'] / \
    df_train['person_age']

# # 6. Loan Percent of Income per Year of Employment
# df_train['loan_percent_emp_year'] = df_train['loan_percent_income'] / \
#     df_train['person_emp_length']

# 7. Default Risk Index
df_train['default_risk_index'] = (df_train['loan_amnt'] * \
    df_train['loan_int_rate']) / \
    (df_train['person_income'] * \
    df_train['cb_person_cred_hist_length'])

# 8. Employment Stability Indicator
df_train['emp_stability_indicator'] = df_train.apply(lambda row: 1 if \
    row['person_emp_length'] > (row['person_age'] * 0.6) else 0, axis=1)

# 9. Debt to Credit History Ratio
df_train['debt_cred_hist_ratio'] = df_train['loan_amnt'] / \
    df_train['cb_person_cred_hist_length']

# Now repeat the same for df_test
# Ensure the 'loan_amnt', 'person_income', etc. columns are present in df_test

# 1. Income to Loan Ratio
df_test['income_loan_ratio'] = df_test['person_income'] / df_test['loan_amnt']

# 2. Age-Employment Length Ratio
df_test['age_emp_length_ratio'] = df_test['person_emp_length'] / \
    df_test['person_age']

# # 3. Loan to Employment Ratio
# df_test['loan_emp_ratio'] = df_test['loan_amnt'] / \
#     df_test['person_emp_length']

# 4. Loan Interest to Income Ratio
df_test['loan_int_income_ratio'] = df_test['loan_int_rate'] / \
    df_test['person_income']

# 5. Credit History to Age Ratio
df_test['cred_hist_age_ratio'] = df_test['cb_person_cred_hist_length'] / \
    df_test['person_age']

# # 6. Loan Percent of Income per Year of Employment

```



```

# df_test['loan_percent_emp_year'] = df_test['loan_percent_income'] /\
↳df_test['person_emp_length']

# 7. Default Risk Index
df_test['default_risk_index'] = (df_test['loan_amnt'] *
↳df_test['loan_int_rate']) /\
(df_test['person_income'] *
↳df_test['cb_person_cred_hist_length'])

# 8. Employment Stability Indicator
df_test['emp_stability_indicator'] = df_test.apply(lambda row: 1 if
↳row['person_emp_length'] > (row['person_age'] * 0.6) else 0, axis=1)

# 9. Debt to Credit History Ratio
df_test['debt_cred_hist_ratio'] = df_test['loan_amnt'] /\
↳df_test['cb_person_cred_hist_length']

# Display the updated df_test dataframe with new features
df_test

```

```

[290]:
      id  person_age  person_income  person_home_ownership \
0    58645         23         69000                RENT
1    58646         26         96000                MORTGAGE
2    58647         26         30000                RENT
3    58648         33         50000                RENT
4    58649         26        102000                MORTGAGE
...    ...         ...         ...         ...
39093  97738         22         31200                MORTGAGE
39094  97739         22         48000                MORTGAGE
39095  97740         51         60000                MORTGAGE
39096  97741         22         36000                MORTGAGE
39097  97742         31         45000                RENT

      person_emp_length  loan_intent  loan_grade  loan_amnt \
0                    3.0  HOMEIMPROVEMENT        F    25000
1                    6.0        PERSONAL        C    10000
2                    5.0        VENTURE        E     4000
3                    4.0  DEBTCONSOLIDATION        A     7000
4                    8.0  HOMEIMPROVEMENT        D    15000
...    ...         ...         ...         ...
39093                2.0  DEBTCONSOLIDATION        B     3000
39094                6.0        EDUCATION        A     7000
39095                0.0        PERSONAL        A    15000
39096                4.0        PERSONAL        D    14000
39097                6.0  DEBTCONSOLIDATION        B    19450

      loan_int_rate  loan_percent_income  cb_person_default_on_file \

```

0	15.76	0.36	N
1	12.68	0.10	Y
2	17.19	0.13	Y
3	8.90	0.14	N
4	16.32	0.15	Y
...
39093	10.37	0.10	N
39094	6.03	0.15	N
39095	7.51	0.25	N
39096	15.62	0.39	Y
39097	9.91	0.44	N

	cb_person_cred_hist_length	income_loan_ratio	age_emp_length_ratio \
0	2	2.760000	0.130435
1	4	9.600000	0.230769
2	2	7.500000	0.192308
3	7	7.142857	0.121212
4	4	6.800000	0.307692
...
39093	4	10.400000	0.090909
39094	3	6.857143	0.272727
39095	25	4.000000	0.000000
39096	4	2.571429	0.181818
39097	9	2.313625	0.193548

	loan_int_income_ratio	cred_hist_age_ratio	default_risk_index \
0	0.000228	0.086957	2.855072
1	0.000132	0.153846	0.330208
2	0.000573	0.076923	1.146000
3	0.000178	0.212121	0.178000
4	0.000160	0.153846	0.600000
...
39093	0.000332	0.181818	0.249279
39094	0.000126	0.136364	0.293125
39095	0.000125	0.490196	0.075100
39096	0.000434	0.181818	1.518611
39097	0.000220	0.290323	0.475925

	emp_stability_indicator	debt_cred_hist_ratio
0	0	12500.000000
1	0	2500.000000
2	0	2000.000000
3	0	1000.000000
4	0	3750.000000
...
39093	0	750.000000
39094	0	2333.333333

```

39095          0          600.000000
39096          0          3500.000000
39097          0          2161.111111

```

[39098 rows x 19 columns]

```

[291]: df_train.drop(columns=['id'], inplace=True)
df_test_ids = df_test['id']
df_test.drop(columns=['id'], inplace=True)

```

```

[292]: # One-Hot Encoding
df_train = pd.get_dummies(df_train, columns=['person_home_ownership',
↳ 'loan_intent', 'loan_grade', 'cb_person_default_on_file'], drop_first=True)
df_test = pd.get_dummies(df_test, columns=['person_home_ownership',
↳ 'loan_intent', 'loan_grade', 'cb_person_default_on_file'], drop_first=True)

```

```

[293]: df_train

```

```

[293]:      person_age  person_income  person_emp_length  loan_amnt  loan_int_rate  \
0              37           35000              0.0         6000         11.49
1              22           56000              6.0         4000         13.35
2              29           28800              8.0         6000          8.90
3              30           70000             14.0        12000         11.11
4              22           60000              2.0         6000          6.92
...          ...          ...          ...          ...          ...
58640          34          120000              5.0        25000         15.95
58641          28           28800              0.0        10000         12.73
58642          23           44000              7.0         6800         16.00
58643          22           30000              2.0         5000          8.90
58644          31           75000              2.0        15000         11.11

```

```

      loan_percent_income  cb_person_cred_hist_length  loan_status  \
0              0.17              14              0
1              0.07              2              0
2              0.21             10              0
3              0.17              5              0
4              0.10              3              0
...          ...          ...          ...
58640          0.21             10              0
58641          0.35              8              1
58642          0.15              2              1
58643          0.17              3              0
58644          0.20              5              0

```

```

      income_loan_ratio  age_emp_length_ratio  ...  loan_intent_MEDICAL  \
0              5.833333              0.000000  ...              False
1              14.000000              0.272727  ...              True

```

2	4.800000	0.275862	...	False
3	5.833333	0.466667	...	False
4	10.000000	0.090909	...	True
...
58640	4.800000	0.147059	...	False
58641	2.880000	0.000000	...	True
58642	6.470588	0.304348	...	False
58643	6.000000	0.090909	...	False
58644	5.000000	0.064516	...	False

	loan_intent_PERSONAL	loan_intent_VENTURE	loan_grade_B	loan_grade_C	\
0	False	False	True	False	
1	False	False	False	True	
2	True	False	False	False	
3	False	True	True	False	
4	False	False	False	False	
...	
58640	False	False	False	False	
58641	False	False	False	True	
58642	False	False	False	False	
58643	False	False	False	False	
58644	False	True	True	False	

	loan_grade_D	loan_grade_E	loan_grade_F	loan_grade_G	\
0	False	False	False	False	
1	False	False	False	False	
2	False	False	False	False	
3	False	False	False	False	
4	False	False	False	False	
...	
58640	True	False	False	False	
58641	False	False	False	False	
58642	True	False	False	False	
58643	False	False	False	False	
58644	False	False	False	False	

	cb_person_default_on_file_Y
0	False
1	False
2	False
3	False
4	False
...	...
58640	True
58641	False
58642	False
58643	False

58644

False

[58645 rows x 30 columns]

```
[294]: df_train[df_train.select_dtypes(include='bool').columns] = df_train.  
        ↪select_dtypes(include='bool').astype(int)  
df_test[df_test.select_dtypes(include='bool').columns] = df_test.  
        ↪select_dtypes(include='bool').astype(int)
```

```
[295]: df_train
```

```
[295]:
```

	person_age	person_income	person_emp_length	loan_amnt	loan_int_rate	\
0	37	35000	0.0	6000	11.49	
1	22	56000	6.0	4000	13.35	
2	29	28800	8.0	6000	8.90	
3	30	70000	14.0	12000	11.11	
4	22	60000	2.0	6000	6.92	
...	
58640	34	120000	5.0	25000	15.95	
58641	28	28800	0.0	10000	12.73	
58642	23	44000	7.0	6800	16.00	
58643	22	30000	2.0	5000	8.90	
58644	31	75000	2.0	15000	11.11	

	loan_percent_income	cb_person_cred_hist_length	loan_status	\
0	0.17	14	0	
1	0.07	2	0	
2	0.21	10	0	
3	0.17	5	0	
4	0.10	3	0	
...	
58640	0.21	10	0	
58641	0.35	8	1	
58642	0.15	2	1	
58643	0.17	3	0	
58644	0.20	5	0	

	income_loan_ratio	age_emp_length_ratio	...	loan_intent_MEDICAL	\
0	5.833333	0.000000	...	0	
1	14.000000	0.272727	...	1	
2	4.800000	0.275862	...	0	
3	5.833333	0.466667	...	0	
4	10.000000	0.090909	...	1	
...	
58640	4.800000	0.147059	...	0	
58641	2.880000	0.000000	...	1	
58642	6.470588	0.304348	...	0	

58643	6.000000	0.090909	...	0
58644	5.000000	0.064516	...	0

	loan_intent_PERSONAL	loan_intent_VENTURE	loan_grade_B	loan_grade_C	\
0	0	0	1	0	
1	0	0	0	1	
2	1	0	0	0	
3	0	1	1	0	
4	0	0	0	0	
...	
58640	0	0	0	0	
58641	0	0	0	1	
58642	0	0	0	0	
58643	0	0	0	0	
58644	0	1	1	0	

	loan_grade_D	loan_grade_E	loan_grade_F	loan_grade_G	\
0	0	0	0	0	
1	0	0	0	0	
2	0	0	0	0	
3	0	0	0	0	
4	0	0	0	0	
...	
58640	1	0	0	0	
58641	0	0	0	0	
58642	1	0	0	0	
58643	0	0	0	0	
58644	0	0	0	0	

	cb_person_default_on_file_Y
0	0
1	0
2	0
3	0
4	0
...	...
58640	1
58641	0
58642	0
58643	0
58644	0

[58645 rows x 30 columns]

```
[296]: numerical_cols_train = df_train.select_dtypes(include=['int64', 'float64']).
        ↪columns.tolist()
        print("Numerical columns in df_train:")
```

```
print(numerical_cols_train)
```

Numerical columns in df_train:

```
['person_age', 'person_income', 'person_emp_length', 'loan_amnt',  
'loan_int_rate', 'loan_percent_income', 'cb_person_cred_hist_length',  
'loan_status', 'income_loan_ratio', 'age_emp_length_ratio',  
'loan_int_income_ratio', 'cred_hist_age_ratio', 'default_risk_index',  
'emp_stability_indicator', 'debt_cred_hist_ratio']
```

```
[297]: # Check for NaN values in df_train and df_test  
print("NaN values in df_train:")  
print(df_train.isna().sum())  
  
print("NaN values in df_test:")  
print(df_test.isna().sum())  
  
# Check for infinite values in df_train and df_test  
print("Infinite values in df_train:")  
print(np.isinf(df_train).sum())  
  
print("Infinite values in df_test:")  
print(np.isinf(df_test).sum())
```

NaN values in df_train:

person_age	0
person_income	0
person_emp_length	0
loan_amnt	0
loan_int_rate	0
loan_percent_income	0
cb_person_cred_hist_length	0
loan_status	0
income_loan_ratio	0
age_emp_length_ratio	0
loan_int_income_ratio	0
cred_hist_age_ratio	0
default_risk_index	0
emp_stability_indicator	0
debt_cred_hist_ratio	0
person_home_ownership_OTHER	0
person_home_ownership_OWN	0
person_home_ownership_RENT	0
loan_intent_EDUCATION	0
loan_intent_HOMEIMPROVEMENT	0
loan_intent_MEDICAL	0
loan_intent_PERSONAL	0
loan_intent_VENTURE	0
loan_grade_B	0

loan_grade_C	0
loan_grade_D	0
loan_grade_E	0
loan_grade_F	0
loan_grade_G	0
cb_person_default_on_file_Y	0
dtype: int64	
NaN values in df_test:	
person_age	0
person_income	0
person_emp_length	0
loan_amnt	0
loan_int_rate	0
loan_percent_income	0
cb_person_cred_hist_length	0
income_loan_ratio	0
age_emp_length_ratio	0
loan_int_income_ratio	0
cred_hist_age_ratio	0
default_risk_index	0
emp_stability_indicator	0
debt_cred_hist_ratio	0
person_home_ownership_OTHER	0
person_home_ownership_OWN	0
person_home_ownership_RENT	0
loan_intent_EDUCATION	0
loan_intent_HOMEIMPROVEMENT	0
loan_intent_MEDICAL	0
loan_intent_PERSONAL	0
loan_intent_VENTURE	0
loan_grade_B	0
loan_grade_C	0
loan_grade_D	0
loan_grade_E	0
loan_grade_F	0
loan_grade_G	0
cb_person_default_on_file_Y	0
dtype: int64	
Infinite values in df_train:	
person_age	0
person_income	0
person_emp_length	0
loan_amnt	0
loan_int_rate	0
loan_percent_income	0
cb_person_cred_hist_length	0
loan_status	0
income_loan_ratio	0

age_emp_length_ratio	0
loan_int_income_ratio	0
cred_hist_age_ratio	0
default_risk_index	0
emp_stability_indicator	0
debt_cred_hist_ratio	0
person_home_ownership_OTHER	0
person_home_ownership_OWN	0
person_home_ownership_RENT	0
loan_intent_EDUCATION	0
loan_intent_HOMEIMPROVEMENT	0
loan_intent_MEDICAL	0
loan_intent_PERSONAL	0
loan_intent_VENTURE	0
loan_grade_B	0
loan_grade_C	0
loan_grade_D	0
loan_grade_E	0
loan_grade_F	0
loan_grade_G	0
cb_person_default_on_file_Y	0
dtype: int64	
Infinite values in df_test:	
person_age	0
person_income	0
person_emp_length	0
loan_amnt	0
loan_int_rate	0
loan_percent_income	0
cb_person_cred_hist_length	0
income_loan_ratio	0
age_emp_length_ratio	0
loan_int_income_ratio	0
cred_hist_age_ratio	0
default_risk_index	0
emp_stability_indicator	0
debt_cred_hist_ratio	0
person_home_ownership_OTHER	0
person_home_ownership_OWN	0
person_home_ownership_RENT	0
loan_intent_EDUCATION	0
loan_intent_HOMEIMPROVEMENT	0
loan_intent_MEDICAL	0
loan_intent_PERSONAL	0
loan_intent_VENTURE	0
loan_grade_B	0
loan_grade_C	0
loan_grade_D	0

```

loan_grade_E      0
loan_grade_F      0
loan_grade_G      0
cb_person_default_on_file_Y  0
dtype: int64

```

```

[298]: # # Replace infinite values with NaN in both training and test data
# df_train.replace([np.inf, -np.inf], np.nan, inplace=True)
# df_test.replace([np.inf, -np.inf], np.nan, inplace=True)

# # Replace NaN values with the median value of each column
# # For df_train
# for column in df_train.columns:
#     if df_train[column].dtype != 'object': # Only apply to numerical columns
#         median_value = df_train[column].median()
#         df_train[column].fillna(median_value, inplace=True)

# # For df_test
# for column in df_test.columns:
#     if df_test[column].dtype != 'object': # Only apply to numerical columns
#         median_value = df_test[column].median()
#         df_test[column].fillna(median_value, inplace=True)

# # After this, you can proceed with your model training and prediction

```

```

[299]: # Train and Test Data Missing Values
print(df_train.isnull().sum())
print(df_test.isnull().sum())

```

```

person_age      0
person_income   0
person_emp_length  0
loan_amnt       0
loan_int_rate   0
loan_percent_income  0
cb_person_cred_hist_length  0
loan_status     0
income_loan_ratio  0
age_emp_length_ratio  0
loan_int_income_ratio  0
cred_hist_age_ratio  0
default_risk_index  0
emp_stability_indicator  0
debt_cred_hist_ratio  0
person_home_ownership_OTHER  0
person_home_ownership_OWN  0
person_home_ownership_RENT  0
loan_intent_EDUCATION  0

```

loan_intent_HOMEIMPROVEMENT	0
loan_intent_MEDICAL	0
loan_intent_PERSONAL	0
loan_intent_VENTURE	0
loan_grade_B	0
loan_grade_C	0
loan_grade_D	0
loan_grade_E	0
loan_grade_F	0
loan_grade_G	0
cb_person_default_on_file_Y	0
dtype: int64	
person_age	0
person_income	0
person_emp_length	0
loan_amnt	0
loan_int_rate	0
loan_percent_income	0
cb_person_cred_hist_length	0
income_loan_ratio	0
age_emp_length_ratio	0
loan_int_income_ratio	0
cred_hist_age_ratio	0
default_risk_index	0
emp_stability_indicator	0
debt_cred_hist_ratio	0
person_home_ownership_OTHER	0
person_home_ownership_OWN	0
person_home_ownership_RENT	0
loan_intent_EDUCATION	0
loan_intent_HOMEIMPROVEMENT	0
loan_intent_MEDICAL	0
loan_intent_PERSONAL	0
loan_intent_VENTURE	0
loan_grade_B	0
loan_grade_C	0
loan_grade_D	0
loan_grade_E	0
loan_grade_F	0
loan_grade_G	0
cb_person_default_on_file_Y	0
dtype: int64	

```
[300]: # Rescaling the relevant numerical features in both training and test sets
scaler = MinMaxScaler()
```

```

num_vars_train = ['person_age', 'person_income', 'person_emp_length',
↳ 'loan_amnt', 'loan_int_rate', 'loan_percent_income',
↳ 'cb_person_cred_hist_length', 'income_loan_ratio', 'age_emp_length_ratio',
↳ 'loan_int_income_ratio', 'cred_hist_age_ratio', 'default_risk_index',
↳ 'emp_stability_indicator', 'debt_cred_hist_ratio']
num_vars_test = ['person_age', 'person_income', 'person_emp_length',
↳ 'loan_amnt', 'loan_int_rate', 'loan_percent_income',
↳ 'cb_person_cred_hist_length', 'income_loan_ratio', 'age_emp_length_ratio',
↳ 'loan_int_income_ratio', 'cred_hist_age_ratio', 'default_risk_index',
↳ 'emp_stability_indicator', 'debt_cred_hist_ratio']
df_train[num_vars_train] = scaler.fit_transform(df_train[num_vars_train])
df_test[num_vars_test] = scaler.transform(df_test[num_vars_test])
from sklearn.preprocessing import StandardScaler

# # Initialize the scaler
# scaler = StandardScaler()

# # List of numerical columns for training and testing sets
# num_vars_train = ['person_age', 'person_income', 'person_emp_length',
↳ 'loan_amnt', 'loan_int_rate', 'loan_percent_income',
↳ 'cb_person_cred_hist_length', 'income_loan_ratio', 'age_emp_length_ratio',
↳ 'loan_int_income_ratio', 'cred_hist_age_ratio', 'default_risk_index',
↳ 'emp_stability_indicator', 'debt_cred_hist_ratio']
# num_vars_test = ['person_age', 'person_income', 'person_emp_length',
↳ 'loan_amnt', 'loan_int_rate', 'loan_percent_income',
↳ 'cb_person_cred_hist_length', 'income_loan_ratio', 'age_emp_length_ratio',
↳ 'loan_int_income_ratio', 'cred_hist_age_ratio', 'default_risk_index',
↳ 'emp_stability_indicator', 'debt_cred_hist_ratio']

# # Standardize the numerical columns in both training and test sets
# df_train[num_vars_train] = scaler.fit_transform(df_train[num_vars_train])
# df_test[num_vars_test] = scaler.transform(df_test[num_vars_test])

# # Display the first few rows to confirm the changes
# df_train.head()

```

```

[301]: # Train and Test Data Missing Values
print(df_train.isnull().sum())
print(df_test.isnull().sum())

```

```

person_age          0
person_income       0
person_emp_length   0
loan_amnt           0
loan_int_rate       0
loan_percent_income 0
cb_person_cred_hist_length 0

```

loan_status	0
income_loan_ratio	0
age_emp_length_ratio	0
loan_int_income_ratio	0
cred_hist_age_ratio	0
default_risk_index	0
emp_stability_indicator	0
debt_cred_hist_ratio	0
person_home_ownership_OTHER	0
person_home_ownership_OWN	0
person_home_ownership_RENT	0
loan_intent_EDUCATION	0
loan_intent_HOMEIMPROVEMENT	0
loan_intent_MEDICAL	0
loan_intent_PERSONAL	0
loan_intent_VENTURE	0
loan_grade_B	0
loan_grade_C	0
loan_grade_D	0
loan_grade_E	0
loan_grade_F	0
loan_grade_G	0
cb_person_default_on_file_Y	0
dtype: int64	
person_age	0
person_income	0
person_emp_length	0
loan_amnt	0
loan_int_rate	0
loan_percent_income	0
cb_person_cred_hist_length	0
income_loan_ratio	0
age_emp_length_ratio	0
loan_int_income_ratio	0
cred_hist_age_ratio	0
default_risk_index	0
emp_stability_indicator	0
debt_cred_hist_ratio	0
person_home_ownership_OTHER	0
person_home_ownership_OWN	0
person_home_ownership_RENT	0
loan_intent_EDUCATION	0
loan_intent_HOMEIMPROVEMENT	0
loan_intent_MEDICAL	0
loan_intent_PERSONAL	0
loan_intent_VENTURE	0
loan_grade_B	0
loan_grade_C	0

```

loan_grade_D          0
loan_grade_E          0
loan_grade_F          0
loan_grade_G          0
cb_person_default_on_file_Y  0
dtype: int64

```

[302]: `df_train.info()`

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 58645 entries, 0 to 58644
Data columns (total 30 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   person_age                            58645 non-null  float64
1   person_income                         58645 non-null  float64
2   person_emp_length                    58645 non-null  float64
3   loan_amnt                            58645 non-null  float64
4   loan_int_rate                        58645 non-null  float64
5   loan_percent_income                 58645 non-null  float64
6   cb_person_cred_hist_length          58645 non-null  float64
7   loan_status                          58645 non-null  int64
8   income_loan_ratio                   58645 non-null  float64
9   age_emp_length_ratio                58645 non-null  float64
10  loan_int_income_ratio                58645 non-null  float64
11  cred_hist_age_ratio                 58645 non-null  float64
12  default_risk_index                  58645 non-null  float64
13  emp_stability_indicator              58645 non-null  float64
14  debt_cred_hist_ratio                 58645 non-null  float64
15  person_home_ownership_OTHER          58645 non-null  int32
16  person_home_ownership_OWNS           58645 non-null  int32
17  person_home_ownership_RENT           58645 non-null  int32
18  loan_intent_EDUCATION                 58645 non-null  int32
19  loan_intent_HOMEIMPROVEMENT           58645 non-null  int32
20  loan_intent_MEDICAL                   58645 non-null  int32
21  loan_intent_PERSONAL                   58645 non-null  int32
22  loan_intent_VENTURE                   58645 non-null  int32
23  loan_grade_B                         58645 non-null  int32
24  loan_grade_C                         58645 non-null  int32
25  loan_grade_D                         58645 non-null  int32
26  loan_grade_E                         58645 non-null  int32
27  loan_grade_F                         58645 non-null  int32
28  loan_grade_G                         58645 non-null  int32
29  cb_person_default_on_file_Y           58645 non-null  int32
dtypes: float64(14), int32(15), int64(1)
memory usage: 10.1 MB

```

```
[303]: df_test.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 39098 entries, 0 to 39097
Data columns (total 29 columns):
 #   Column                                  Non-Null Count  Dtype
---  -
 0   person_age                             39098 non-null  float64
 1   person_income                           39098 non-null  float64
 2   person_emp_length                       39098 non-null  float64
 3   loan_amnt                               39098 non-null  float64
 4   loan_int_rate                           39098 non-null  float64
 5   loan_percent_income                     39098 non-null  float64
 6   cb_person_cred_hist_length              39098 non-null  float64
 7   income_loan_ratio                       39098 non-null  float64
 8   age_emp_length_ratio                    39098 non-null  float64
 9   loan_int_income_ratio                    39098 non-null  float64
10   cred_hist_age_ratio                     39098 non-null  float64
11   default_risk_index                      39098 non-null  float64
12   emp_stability_indicator                  39098 non-null  float64
13   debt_cred_hist_ratio                    39098 non-null  float64
14   person_home_ownership_OTHER              39098 non-null  int32
15   person_home_ownership_OWEN               39098 non-null  int32
16   person_home_ownership_RENT               39098 non-null  int32
17   loan_intent_EDUCATION                    39098 non-null  int32
18   loan_intent_HOMEIMPROVEMENT              39098 non-null  int32
19   loan_intent_MEDICAL                      39098 non-null  int32
20   loan_intent_PERSONAL                     39098 non-null  int32
21   loan_intent_VENTURE                      39098 non-null  int32
22   loan_grade_B                             39098 non-null  int32
23   loan_grade_C                             39098 non-null  int32
24   loan_grade_D                             39098 non-null  int32
25   loan_grade_E                             39098 non-null  int32
26   loan_grade_F                             39098 non-null  int32
27   loan_grade_G                             39098 non-null  int32
28   cb_person_default_on_file_Y              39098 non-null  int32
dtypes: float64(14), int32(15)
memory usage: 6.4 MB
```

```
[304]: # Prepare the training and validation data
X_train = df_train.drop('loan_status', axis=1) # Features
y_train = df_train['loan_status'] # Target
```

```
[ ]:
```

```
[305]: X_train, X_test, y_train, y_test = train_test_split(X_train, y_train,
↳ test_size=0.2, random_state=42)
```

```
[306]: # #Handle Imbalance using SMOTE for oversampling the minority class
# smote = SMOTE(random_state=42)
# X_resampled, y_resampled = smote.fit_resample(X_train, y_train)
```

```
[307]: X_resampled
```

```
[307]:      person_age  person_income  person_emp_length  loan_amnt  loan_int_rate  \
0      0.155340      0.024159      0.000000      0.194203      0.473596
1      0.067961      0.020572      0.024390      0.130435      0.350562
2      0.048544      0.034708      0.000000      0.217391      0.138202
3      0.184466      0.037346      0.016260      0.420290      0.062921
4      0.019417      0.021838      0.040650      0.101449      0.125843
...      ...      ...      ...      ...      ...
80411      0.044037      0.032021      0.018776      0.663220      0.241431
80412      0.031521      0.050436      0.035475      0.289970      0.548212
80413      0.116505      0.047708      0.069322      0.420290      0.628559
80414      0.097087      0.026870      0.008384      0.130989      0.492537
80415      0.114145      0.029215      0.014576      0.565217      0.311050

      loan_percent_income  cb_person_cred_hist_length  income_loan_ratio  \
0      0.168675      0.464286      0.005231
1      0.144578      0.250000      0.006570
2      0.132530      0.035714      0.006657
3      0.240964      0.321429      0.003696
4      0.108434      0.000000      0.008750
...      ...      ...      ...
80411      0.455104      0.000000      0.001932
80412      0.128152      0.013452      0.007475
80413      0.192771      0.250000      0.004730
80414      0.108622      0.142300      0.008425
80415      0.400086      0.285714      0.002100

      age_emp_length_ratio  loan_int_income_ratio  ...  loan_intent_MEDICAL  \
0      0.000000      0.104558  ...      0
1      0.018970      0.101838  ...      0
2      0.000000      0.041502  ...      0
3      0.008755      0.031773  ...      0
4      0.038803      0.062752  ...      1
...      ...      ...      ...
80411      0.016348      0.056861  ...      0
80412      0.031041      0.058423  ...      0
80413      0.045493      0.065648  ...      0
80414      0.005869      0.097019  ...      0
80415      0.009514      0.068835  ...      0

      loan_intent_PERSONAL  loan_intent_VENTURE  loan_grade_B  loan_grade_C  \
0      0      0      0      1
```


1		0		0		0		1
2		0		0		0		0
3		0		0		0		0
4		0		0		0		0
...	
80411		0		0		1		0
80412		0		0		0		0
80413		0		0		0		0
80414		0		0		0		1
80415		0		0		1		0

	loan_grade_D	loan_grade_E	loan_grade_F	loan_grade_G	\
0	0	0	0	0	
1	0	0	0	0	
2	0	0	0	0	
3	0	0	0	0	
4	0	0	0	0	
...	
80411	0	0	0	0	
80412	1	0	0	0	
80413	1	0	0	0	
80414	0	0	0	0	
80415	0	0	0	0	

	cb_person_default_on_file_Y
0	1
1	1
2	0
3	0
4	0
...	...
80411	0
80412	1
80413	1
80414	0
80415	0

[80416 rows x 29 columns]

```
[308]: ## Initialize the Random Forest model
# rf_model = RandomForestClassifier(random_state=42)

## Train the model on the resampled data
# rf_model.fit(X_resampled, y_resampled)

## Predict on the validation set
# y_pred = rf_model.predict(X_test)
```

```

# # Calculate the accuracy
# accuracy = accuracy_score(y_test, y_pred)
# print(f"Accuracy: {accuracy:.2f}")

# # Print the classification report for more detailed performance analysis
# print(classification_report(y_test, y_pred))

```

Accuracy: 0.94

	precision	recall	f1-score	support
0	0.96	0.97	0.97	10087
1	0.81	0.75	0.78	1642
accuracy			0.94	11729
macro avg	0.89	0.86	0.87	11729
weighted avg	0.94	0.94	0.94	11729

```

[309]: # test_predictions = rf_model.predict(df_test)

# # Create a DataFrame with test IDs and predictions
# output = pd.DataFrame({
#     'id': df_test_ids, # Use the test IDs from earlier
#     'loan_status': test_predictions # Predictions from the model
# })

# # Export the results to an Excel file
# output.to_excel('loan_status_predictions.xlsx', index=False)

```

```

[310]: from catboost import CatBoostClassifier
catboost_model = CatBoostClassifier(random_state=42, verbose=0)

# Train the model on the resampled data
catboost_model.fit(X_resampled, y_resampled)

# Predict on the validation set (X_test)
y_pred = catboost_model.predict(X_test)

# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f"CatBoost Accuracy: {accuracy:.2f}")

# Print the classification report for detailed performance analysis
print(classification_report(y_test, y_pred))

# Optionally, you can save the model's predictions on the test dataset

```

```
test_predictions = catboost_model.predict(df_test)

# Create a DataFrame with test IDs and predictions
output = pd.DataFrame({
    'id': df_test_ids,
    'loan_status': test_predictions
})

# Export the predictions to an Excel file
output.to_excel('catboost_loan_status_predictions.xlsx', index=False)
```

CatBoost Accuracy: 0.95

	precision	recall	f1-score	support
0	0.96	0.99	0.97	10087
1	0.90	0.74	0.81	1642
accuracy			0.95	11729
macro avg	0.93	0.87	0.89	11729
weighted avg	0.95	0.95	0.95	11729

```
[311]: print(f"Length of df_test_ids: {len(df_test_ids)}")
print(f"Length of test_predictions: {len(test_predictions)}")
```

Length of df_test_ids: 39098

Length of test_predictions: 39098

```
[312]: import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
from sklearn.metrics import accuracy_score, classification_report

# Define the neural network model
model = Sequential()

# Input layer
model.add(Dense(128, activation='relu', input_shape=(X_resampled.shape[1],)))
    ↪ # Adjust input shape based on your feature set
model.add(Dropout(0.3)) # Dropout layer to prevent overfitting

# Hidden layers
model.add(Dense(64, activation='relu'))
model.add(Dropout(0.3))

model.add(Dense(32, activation='relu'))
```

```

model.add(Dropout(0.3))

# Output layer (1 neuron with sigmoid activation for binary classification)
model.add(Dense(1, activation='sigmoid'))

# Compile the model
model.compile(optimizer='adam', loss='binary_crossentropy',
              metrics=['accuracy'])

# Train the model
history = model.fit(X_resampled, y_resampled, epochs=50, batch_size=32,
                  validation_split=0.2)

# Predict on the test set
y_pred = (model.predict(X_test) > 0.5).astype("int32")

# Calculate the accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.2f}")

# Print the classification report for more detailed performance analysis
print(classification_report(y_test, y_pred))

```

Epoch 1/50

C:\Users\Dell\anaconda3\Lib\site-packages\keras\src\layers\core\dense.py:87:
 UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When
 using Sequential models, prefer using an `Input(shape)` object as the first
 layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

2011/2011 5s 2ms/step -
 accuracy: 0.8056 - loss: 0.4359 - val_accuracy: 0.7586 - val_loss: 0.4776

Epoch 2/50

2011/2011 3s 2ms/step -
 accuracy: 0.8739 - loss: 0.3253 - val_accuracy: 0.7603 - val_loss: 0.4225

Epoch 3/50

2011/2011 3s 2ms/step -
 accuracy: 0.8772 - loss: 0.3139 - val_accuracy: 0.7690 - val_loss: 0.4131

Epoch 4/50

2011/2011 3s 1ms/step -
 accuracy: 0.8819 - loss: 0.3077 - val_accuracy: 0.7605 - val_loss: 0.4052

Epoch 5/50

2011/2011 3s 2ms/step -
 accuracy: 0.8799 - loss: 0.3041 - val_accuracy: 0.7809 - val_loss: 0.3820

Epoch 6/50

2011/2011 3s 2ms/step -
 accuracy: 0.8810 - loss: 0.3046 - val_accuracy: 0.7507 - val_loss: 0.4524

Epoch 7/50
 2011/2011 3s 2ms/step -
 accuracy: 0.8808 - loss: 0.2988 - val_accuracy: 0.7636 - val_loss: 0.4272
 Epoch 8/50
 2011/2011 3s 2ms/step -
 accuracy: 0.8812 - loss: 0.3005 - val_accuracy: 0.7518 - val_loss: 0.4329
 Epoch 9/50
 2011/2011 3s 2ms/step -
 accuracy: 0.8811 - loss: 0.2947 - val_accuracy: 0.7476 - val_loss: 0.4335
 Epoch 10/50
 2011/2011 3s 2ms/step -
 accuracy: 0.8835 - loss: 0.2918 - val_accuracy: 0.7294 - val_loss: 0.5167
 Epoch 11/50
 2011/2011 3s 2ms/step -
 accuracy: 0.8817 - loss: 0.2969 - val_accuracy: 0.7592 - val_loss: 0.4305
 Epoch 12/50
 2011/2011 3s 2ms/step -
 accuracy: 0.8848 - loss: 0.2883 - val_accuracy: 0.7711 - val_loss: 0.3941
 Epoch 13/50
 2011/2011 3s 2ms/step -
 accuracy: 0.8847 - loss: 0.2874 - val_accuracy: 0.7643 - val_loss: 0.3893
 Epoch 14/50
 2011/2011 3s 2ms/step -
 accuracy: 0.8883 - loss: 0.2836 - val_accuracy: 0.7583 - val_loss: 0.4092
 Epoch 15/50
 2011/2011 3s 2ms/step -
 accuracy: 0.8854 - loss: 0.2889 - val_accuracy: 0.7429 - val_loss: 0.4473
 Epoch 16/50
 2011/2011 3s 2ms/step -
 accuracy: 0.8848 - loss: 0.2852 - val_accuracy: 0.7507 - val_loss: 0.4334
 Epoch 17/50
 2011/2011 3s 2ms/step -
 accuracy: 0.8854 - loss: 0.2849 - val_accuracy: 0.7562 - val_loss: 0.4379
 Epoch 18/50
 2011/2011 3s 2ms/step -
 accuracy: 0.8850 - loss: 0.2868 - val_accuracy: 0.7598 - val_loss: 0.4415
 Epoch 19/50
 2011/2011 3s 1ms/step -
 accuracy: 0.8846 - loss: 0.2863 - val_accuracy: 0.7423 - val_loss: 0.4679
 Epoch 20/50
 2011/2011 3s 2ms/step -
 accuracy: 0.8856 - loss: 0.2838 - val_accuracy: 0.7551 - val_loss: 0.4316
 Epoch 21/50
 2011/2011 3s 2ms/step -
 accuracy: 0.8870 - loss: 0.2827 - val_accuracy: 0.7683 - val_loss: 0.4111
 Epoch 22/50
 2011/2011 3s 2ms/step -
 accuracy: 0.8877 - loss: 0.2805 - val_accuracy: 0.7542 - val_loss: 0.3992

Epoch 23/50
 2011/2011 3s 1ms/step -
 accuracy: 0.8874 - loss: 0.2806 - val_accuracy: 0.7769 - val_loss: 0.3755
 Epoch 24/50
 2011/2011 3s 1ms/step -
 accuracy: 0.8856 - loss: 0.2819 - val_accuracy: 0.7580 - val_loss: 0.4143
 Epoch 25/50
 2011/2011 3s 2ms/step -
 accuracy: 0.8873 - loss: 0.2801 - val_accuracy: 0.7638 - val_loss: 0.4015
 Epoch 26/50
 2011/2011 3s 1ms/step -
 accuracy: 0.8891 - loss: 0.2773 - val_accuracy: 0.7622 - val_loss: 0.3833
 Epoch 27/50
 2011/2011 3s 2ms/step -
 accuracy: 0.8889 - loss: 0.2755 - val_accuracy: 0.7609 - val_loss: 0.4049
 Epoch 28/50
 2011/2011 3s 2ms/step -
 accuracy: 0.8883 - loss: 0.2743 - val_accuracy: 0.7672 - val_loss: 0.3996
 Epoch 29/50
 2011/2011 3s 2ms/step -
 accuracy: 0.8870 - loss: 0.2785 - val_accuracy: 0.7609 - val_loss: 0.4252
 Epoch 30/50
 2011/2011 3s 2ms/step -
 accuracy: 0.8884 - loss: 0.2750 - val_accuracy: 0.7680 - val_loss: 0.3960
 Epoch 31/50
 2011/2011 3s 2ms/step -
 accuracy: 0.8884 - loss: 0.2780 - val_accuracy: 0.7538 - val_loss: 0.4167
 Epoch 32/50
 2011/2011 3s 2ms/step -
 accuracy: 0.8889 - loss: 0.2732 - val_accuracy: 0.7612 - val_loss: 0.4162
 Epoch 33/50
 2011/2011 3s 2ms/step -
 accuracy: 0.8888 - loss: 0.2749 - val_accuracy: 0.7713 - val_loss: 0.4198
 Epoch 34/50
 2011/2011 3s 2ms/step -
 accuracy: 0.8904 - loss: 0.2734 - val_accuracy: 0.7701 - val_loss: 0.3713
 Epoch 35/50
 2011/2011 3s 2ms/step -
 accuracy: 0.8874 - loss: 0.2772 - val_accuracy: 0.7680 - val_loss: 0.3810
 Epoch 36/50
 2011/2011 3s 2ms/step -
 accuracy: 0.8901 - loss: 0.2705 - val_accuracy: 0.7663 - val_loss: 0.3911
 Epoch 37/50
 2011/2011 3s 2ms/step -
 accuracy: 0.8900 - loss: 0.2692 - val_accuracy: 0.7604 - val_loss: 0.4015
 Epoch 38/50
 2011/2011 3s 2ms/step -
 accuracy: 0.8879 - loss: 0.2769 - val_accuracy: 0.7471 - val_loss: 0.4369

Epoch 39/50
 2011/2011 3s 2ms/step -
 accuracy: 0.8903 - loss: 0.2706 - val_accuracy: 0.7716 - val_loss: 0.3895
 Epoch 40/50
 2011/2011 3s 2ms/step -
 accuracy: 0.8895 - loss: 0.2714 - val_accuracy: 0.7671 - val_loss: 0.4066
 Epoch 41/50
 2011/2011 3s 2ms/step -
 accuracy: 0.8895 - loss: 0.2701 - val_accuracy: 0.7614 - val_loss: 0.4101
 Epoch 42/50
 2011/2011 3s 2ms/step -
 accuracy: 0.8905 - loss: 0.2731 - val_accuracy: 0.7714 - val_loss: 0.3864
 Epoch 43/50
 2011/2011 3s 2ms/step -
 accuracy: 0.8891 - loss: 0.2716 - val_accuracy: 0.7652 - val_loss: 0.4146
 Epoch 44/50
 2011/2011 3s 2ms/step -
 accuracy: 0.8908 - loss: 0.2693 - val_accuracy: 0.7675 - val_loss: 0.4197
 Epoch 45/50
 2011/2011 3s 2ms/step -
 accuracy: 0.8879 - loss: 0.2710 - val_accuracy: 0.7678 - val_loss: 0.3858
 Epoch 46/50
 2011/2011 3s 2ms/step -
 accuracy: 0.8877 - loss: 0.2722 - val_accuracy: 0.7696 - val_loss: 0.3880
 Epoch 47/50
 2011/2011 3s 2ms/step -
 accuracy: 0.8894 - loss: 0.2686 - val_accuracy: 0.7693 - val_loss: 0.3935
 Epoch 48/50
 2011/2011 3s 2ms/step -
 accuracy: 0.8891 - loss: 0.2679 - val_accuracy: 0.7822 - val_loss: 0.3704
 Epoch 49/50
 2011/2011 3s 2ms/step -
 accuracy: 0.8900 - loss: 0.2698 - val_accuracy: 0.7702 - val_loss: 0.3769
 Epoch 50/50
 2011/2011 3s 2ms/step -
 accuracy: 0.8891 - loss: 0.2707 - val_accuracy: 0.7687 - val_loss: 0.4140
 367/367 0s 1ms/step
 Accuracy: 0.94

	precision	recall	f1-score	support
0	0.96	0.97	0.96	10087
1	0.78	0.75	0.77	1642
accuracy			0.94	11729
macro avg	0.87	0.86	0.86	11729
weighted avg	0.93	0.94	0.94	11729

```
[313]: # Make predictions on the test set
test_predictions = model.predict(df_test)

# Flatten the predictions array if necessary
test_predictions = test_predictions.flatten()

# Create a DataFrame with test IDs and predictions
output = pd.DataFrame({
    'id': df_test_ids,
    'loan_status': test_predictions
})

# Export the predictions to an Excel file
output.to_excel('catboost_loan_status_predictions.xlsx', index=False)
```

1222/1222

1s 779us/step

[]: